

PostgreSQL

Déploiement de l'opérateur PostgreSQL de Zalando

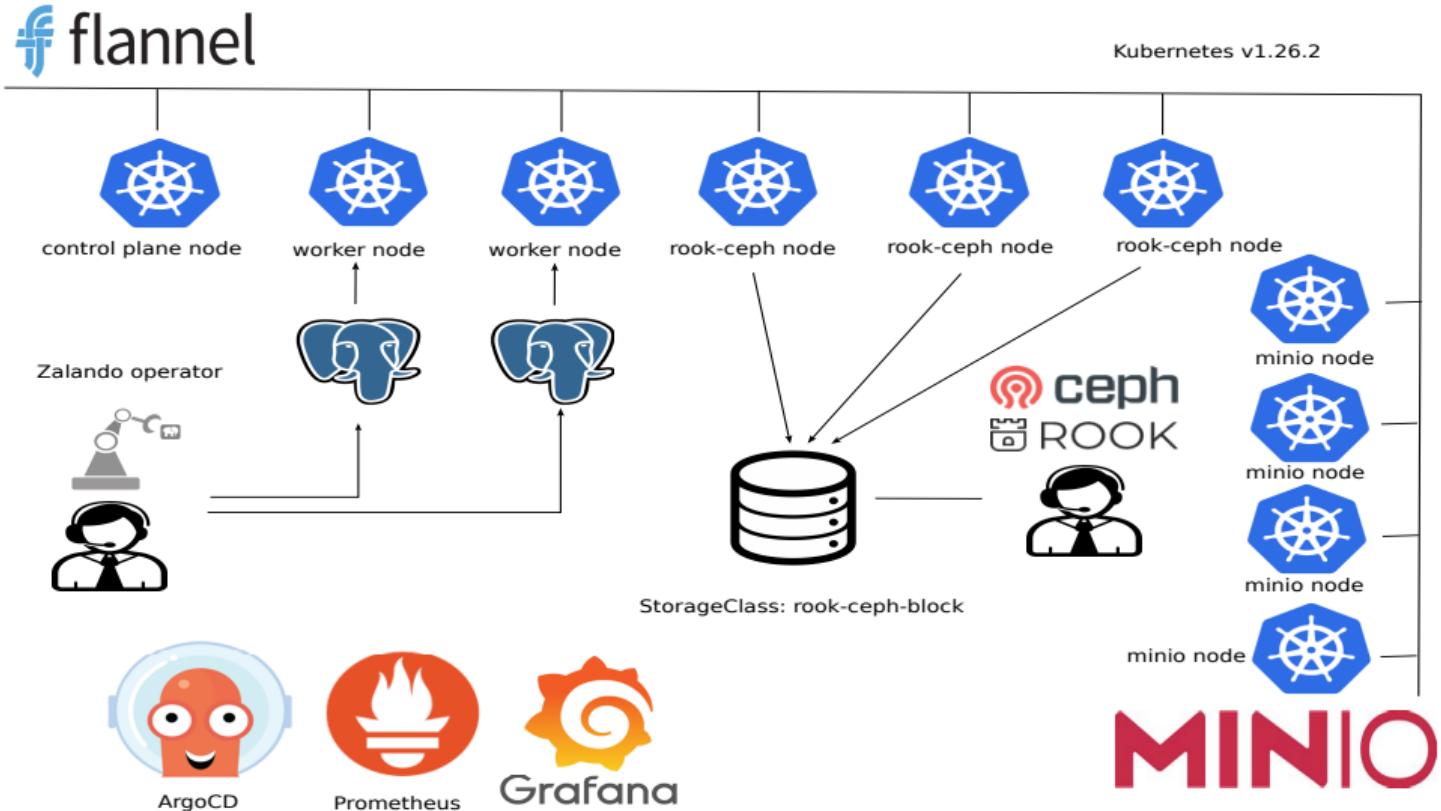
Simon ELBAZ (selbaz@linagora.com / ossa@linagora.com)

27 février 2023



Architecture

Architecture



Installation

Versions utilisées

- OS de déploiement : Debian 11 - Bullseye
- Versions de Kubernetes : 1.26.x

Dimensionnement des serveurs

Dimensionnement du control plane :

- 8 CPU
- 8 Go RAM

Dimensionnement des workers :

- 2 CPU
- 2 Go RAM

Prérequis matériels

- etcd est la base de données clé-valeurs centrale utilisée par Kubernetes
- etcd utilise de manière intensive les disques à disposition
- Pour une stabilité accrue du cluster, il est préférable d'utiliser des disques de type **SSD**

Déploiement du nœud control plane

- Kubernetes s'appuie sur un élément essentiel qui est le *container runtime*.
- La méthode de déploiement du container runtime s'appuie la méthode décrite dans le lien :
<https://docs.docker.com/engine/install/debian/>

Désactivation permanente de la mémoire swap

Le process kubelet ne démarre pas en cas de mémoire swap activée.
Pour désactiver l'utilisation de la swap, merci d'utiliser la commande suivante :

```
swapoff -a
```

Pour persister cet état et faire en sorte que la mémoire swap ne soit pas activée au prochain reboot, supprimer ou mettre en commentaires la ligne suivante dans */etc/fstab* :

```
$ sudo cat /etc/fstab
/dev/mapper/dnumworker1--vg-root /          ext4    errors=remount-ro 0      1
# /boot was on /dev/sda1 during installation
UUID=ddd6fd9d-6ac3-4510-9156-22984bc82b67 /boot      ext2    defaults      0      2
#/dev/mapper/dnumworker1--vg-swap_1 none      swap    sw        0      0
/dev/sr0      /media/cdrom0  udf,iso9660 user,noauto  0      0
```

Installation du runtime container *containerd*

Mise à jour de l'index du paquet *apt* et installation des paquets nécessaires à l'utilisation des dépôts avec le protocole HTTPS :

```
sudo apt-get update  
  
sudo apt-get install \  
  ca-certificates \  
  curl \  
  gnupg
```

Ajout de la clef GPG officielle de Docker

```
sudo install -m 0755 -d /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/debian/gpg \
| sudo gpg --dearmor -o /etc/apt/keyrings/docker.gpg
sudo chmod a+r /etc/apt/keyrings/docker.gpg
```

Ajout du dépôt de Docker

```
echo \  
"deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg] https://download.docker.com/linux/debian \  
"$(. /etc/os-release & echo "$VERSION_CODENAME")" stable" | \  
sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

Installation de Docker Engine

```
sudo apt-get update  
sudo apt-get install docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose-plugin
```

Installation de kubectl, kubeadm et kubelet

```
sudo apt-get install -y ca-certificates curl  
sudo apt-get install -y apt-transport-https  
sudo apt-get update  
curl -fsSL https://packages.cloud.google.com/apt/doc/apt-key.gpg \  
| sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-archive-keyring.gpg  
echo "deb [signed-by=/etc/apt/keyrings/kubernetes-archive-keyring.gpg] \  
https://apt.kubernetes.io/ kubernetes-xenial main" | \  
sudo tee /etc/apt/sources.list.d/kubernetes.list  
  
sudo apt-get update  
sudo apt-get install -y kubectl  
sudo apt-get install -y kubeadm  
sudo apt-get install -y kubelet
```

Activation des modules kernel *overlay* et *br_nfILTER*

```
linagora@debian-cp:/etc/modules-load.d$ cat k8s.conf
overlay
br_nfILTER
linagora@debian-cp:/etc/modules-load.d$ pwd
/etc/modules-load.d
```

Activation des fonctions *bridge/iptables* et *forward* du kernel

```
linagora@debian-cp:/etc/sysctl.d$ cat k8s.conf
inet.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
linagora@debian-cp:/etc/sysctl.d$ pwd
/etc/sysctl.d
```

Paramétrage de containerd

Génération du paramétrage par défaut de containerd :

```
root@debian-cp:~# containerd config default dump > /etc/containerd/config.toml.dmp
```

Modifier la valeur à **true** pour le paramètre **SystemdCgroup** :

```
[plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc.options]
  BinaryName = ""
  CriuImagePath = ""
  CriuPath = ""
  CriuWorkPath = ""
  IoGid = 0
  IoUid = 0
  NoNewKeyring = false
  NoPivotRoot = false
  Root = ""
  ShimCgroup = ""
  SystemdCgroup = true
```

Paramétrage de containerd

Remplacer le paramétrage actuel par le paramétrage modifié :

```
root@debian-cp:~# cp /etc/containerd/config.toml /etc/containerd/config.toml.bak
root@debian-cp:~# cat /etc/containerd/config.toml.dmp > /etc/containerd/config.toml
root@debian-cp:~# systemctl restart containerd
```

Initialisation du cluster Kubernetes

En tant que root, lancer la commande suivante :

```
# kubeadm init --control-plane-endpoint 10.10.10.30 \
--skip-phases=addon/coredns,addon/kube-proxy \
--v=5 \
--pod-network-cidr="10.244.0.0/16"
```

Si les phases *addon/coredns* et *addon/kube-proxy* ne sont pas évitées au 1^{er} lancement de kubeadm, l'erreur suivante est générée :

```
[kubelet-finalize]Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key error execution phase
addon/coredns : unable to fetch CoreDNS current installed version and ConfigMap. : rpc error : code = Unknown desc = malformed header :
missing HTTP content-type To see the stack trace of this error execute with --v=5 or higher
```

Initialisation du cluster Kubernetes

Le résultat de la commande d'init est le suivant :

```
I0315 01:06:38.342010 34405 kubeletfinalize.go:134] [kubelet-finalize] Restarting the
kubelet to enable client certificate rotation
Your Kubernetes control-plane has initialized successfully!
To start using your cluster, you need to run the following as a regular user:
```

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run :

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at :

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

You can now join any number of control-plane nodes by copying certificate authorities and service account keys on each node and then running the following as root :

```
kubeadm join 10.10.10.30:6443 --token 6pia7c.n6u8pbm7yjl6nnr8 \
--discovery-token-ca-cert-hash sha256:f6d45602ea75c7659dc91f661d19e97e6817e2847e4e5d0047880b871317a145 \
--control-plane
```

Then you can join any number of worker nodes by running the following on each as root :

```
kubeadm join 10.10.10.30:6443 --token 6pia7c.n6u8pbm7yjl6nnr8 \
--discovery-token-ca-cert-hash sha256:f6d45602ea75c7659dc91f661d19e97e6817e2847e4e5d0047880b871317a145
```

Paramétrage de *kubectl*

L'utilisation de *kubectl* nécessite l'action suivante :

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Déploiement de l'addon CoreDNS

Comme indiqué précédemment, les addons CoreDNS et Kube-Proxy n'ont pas été déployés au *1^{er}* lancement de kubeadm.

CoreDNS peut maintenant être déployé sans erreur :

```
linagora@debian-cp:~$ sudo kubeadm init phase addon coredns
[addons] Applied essential addon: CoreDNS
```

Déploiement de l'addon Kube-Proxy

```
linagora@debian-cp:~$ sudo kubeadm init phase addon kube-proxy
[addons] Applied essential addon: kube-proxy
```

Choix de la couche réseau - Container Network Interface

Il existe différentes addons Kubernetes implémentant l'interface CNI.
Ces addons sont listés dans l'URL suivante : <https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Pour le POC, l'addon sélectionné est Flannel car il semble être le plus simple et le plus basique des addons CNI.

Déploiement de l'addon *Flannel*

L'addon Flannel s'installe de plusieurs manières (<https://github.com/flannel-io/flannel#deploying-flannel-manually>). La méthode utilisée pour le POC est kubectl :

```
kubectl apply -f https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yml
```

Installation de k9s

Un outil pratique de visualisation d'un cluster kubernetes est : **k9s**
(<https://k9scli.io/>)

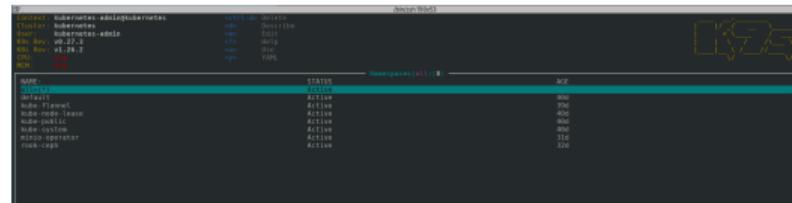
Le lien suivant¹ permet de télécharger l'archive incluant le binaire.

1.

https://github.com/derailed/k9s/releases/download/v0.27.3/k9s_Linux_amd64.tar.gz

Liste des namespaces

```
linagora@debian-cp:~$ kubectl get namespaces  
NAME        STATUS   AGE  
default     Active   40d  
kube-flannel Active   39d  
kube-node-lease Active   40d  
kube-public  Active   40d  
kube-system  Active   40d  
minio-operator Active   32d  
rook-ceph    Active   32d
```

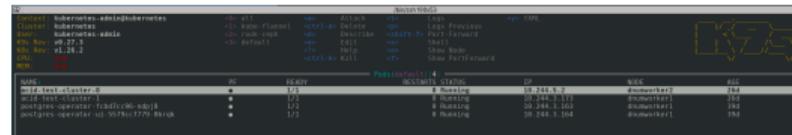


NAME	STATUS	AGE
default	Active	40d
kube-flannel	Active	39d
kube-node-lease	Active	40d
kube-public	Active	40d
kube-system	Active	40d
minio-operator	Active	32d
rook-ceph	Active	32d

Pods du namespace default

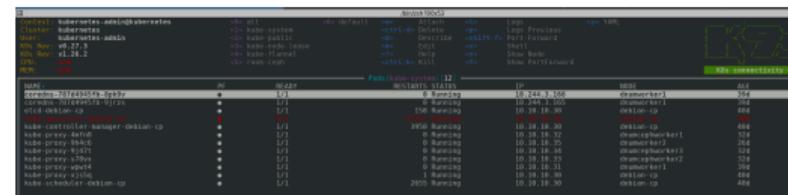
```
linagora@debian-cp:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
acid-test-cluster-0	1/1	Running	0	27d
acid-test-cluster-1	1/1	Running	0	27d
postgres-operator-fcbd7cc96-ndpj8	1/1	Running	0	40d
postgres-operator-ui-5579cc7779-86rqk	1/1	Running	0	40d



Pods du namespace kube-system

```
linagora@debian-cp:~$ kubectl get pods -n kube-system
NAME                      READY   STATUS    RESTARTS   AGE
coredns-787d4945fb-8ph9v   1/1    Running   0          40d
coredns-787d4945fb-9jrzs   1/1    Running   0          40d
etcd-debian-cp              1/1    Running   158        41d
kube-apiserver-debian-cp   0/1    Running   4968 (13m ago) 41d
kube-controller-manager-debian-cp 1/1    Running   4161 (8m26s ago) 41d
kube-proxy-4mf8             1/1    Running   0          33d
kube-proxy-9h4c6             1/1    Running   0          27d
kube-proxy-9j47t             1/1    Running   0          33d
kube-proxy-s78vx             1/1    Running   0          33d
kube-proxy-wpwt4             1/1    Running   0          40d
kube-proxy-xjs5q             1/1    Running   1 (33d ago) 41d
kube-scheduler-debian-cp    1/1    Running   2848 (6m20s ago) 41d
```



Pods du namespace kube-flannel

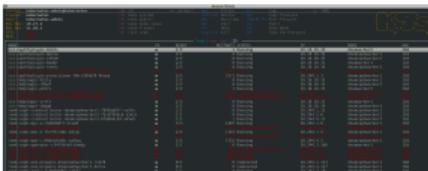
```
linagora@debian-cp:~$ kubectl get pods -n kube-flannel
NAME                READY   STATUS    RESTARTS   AGE
kube-flannel-ds-5nw2j 1/1     Running   0          33d
kube-flannel-ds-5xwsm 1/1     Running   0          40d
kube-flannel-ds-8vkg9  1/1     Running   1 (33d ago) 40d
kube-flannel-ds-pv6ss  1/1     Running   0          27d
kube-flannel-ds-trbz9 1/1     Running   0          33d
kube-flannel-ds-wmzz2  1/1     Running   0          33d
```

The screenshot shows a terminal window with two parts. The top part displays the command output from 'kubectl get pods -n kube-flannel'. The bottom part is a network diagram titled 'Network Diagram' showing nodes and their connections.

NAME	READY	STATUS	RESTARTS	AGE
kube-flannel-ds-5nw2j	1/1	Running	0	33d
kube-flannel-ds-5xwsm	1/1	Running	0	40d
kube-flannel-ds-8vkg9	1/1	Running	1 (33d ago)	40d
kube-flannel-ds-pv6ss	1/1	Running	0	27d
kube-flannel-ds-trbz9	1/1	Running	0	33d
kube-flannel-ds-wmzz2	1/1	Running	0	33d

Pods du namespace rook-ceph

```
linagora@debian-cp:~$ kubectl get pods -n rook-ceph
NAME                                         READY   STATUS    RESTARTS   AGE
csi-cephfsplugin-9nbts                      2/2    Running   1 (27d ago)  27d
csi-cephfsplugin-bpxlw                       2/2    Running   0          33d
csi-cephfsplugin-jd5x8                       2/2    Running   0          33d
csi-cephfsplugin-mddkf                       2/2    Running   0          33d
csi-cephfsplugin-nrmfz                       2/2    Running   0          33d
csi-cephfsplugin-provisioner-84cc595b78-9mml4 5/5    Running   6008 (2m44s ago) 33d
csi-cephfsplugin-provisioner-84cc595b78-9twnq  5/5    Running   2171      33d
csi-rbdplugin-922lq                          2/2    Running   0          33d
csi-rbdplugin-c95w7                          2/2    Running   0          33d
csi-rbdplugin-pk57s                         2/2    Running   1 (27d ago)  27d
csi-rbdplugin-provisioner-6f6b6b8cd6-4c8jd  1/5    CreateContainerError 1344     33d
csi-rbdplugin-provisioner-6f6b6b8cd6-gw6bm   1/5    CreateContainerError 4465     33d
csi-rbdplugin-srtfz                         2/2    Running   0          33d
csi-rbdplugin-v6gqm                         2/2    Running   0          33d
rook-ceph-crashcollector-dnumcephworker1-7845bb8ff-vs9fx 1/1    Running   0          32d
rook-ceph-crashcollector-dnumcephworker2-75cdf95dcd-n5xsz 1/1    Running   0          32d
rook-ceph-crashcollector-dnumcephworker3-6fdb6cd9-x45w5  1/1    Running   1 (8d ago)  32d
rook-ceph-mgr-a-c5db5dff-hvsp9               3/3    Running   1487 (6d6h ago) 33d
rook-ceph-mgr-b-7bfbd88c8b-wh4ww             2/3    CreateContainerError 944      22d
rook-ceph-mon-a-75cf9ccdc-b2jgc              2/2    Running   1163      33d
rook-ceph-mon-b-78d6586d5-qss4z              1/2    CreateContainerError 701 (19d ago) 19d
rook-ceph-mon-c-64dcba86c-w28sg              2/2    Running   1755      33d
rook-ceph-operator-cf4f7df4-6tm6p             1/1    Running   0          32d
rook-ceph-osd-0-579bbdb4d-d6dhrr            1/2    CreateContainerError 484      32d
rook-ceph-osd-1-74698f77fd-6n2mh             1/2    Running   529       32d
rook-ceph-osd-2-5cc486467c-lhm47             1/2    Running   1116 (49m ago) 32d
rook-ceph-osd-prepare-dnumcephworker1-rnk78  0/1    Completed  0          21d
rook-ceph-osd-prepare-dnumcephworker3-42rxv  0/1    Completed  0          21d
rook-ceph-tools-7c4b8bb9b5-pxk67              1/1    Running   0          33d
```



Déploiement du nœud worker

Sur chacun des 2 workers, il est nécessaire de déployer :

- le runtime containerd de Docker
- les commandes kubectl, kubeadm et kubelet
- l'activation des modules kernel overlay et br_netfilter
- l'activation des fonctions bridge/iptables et forward du kernel
- le paramétrage de containerd

Ajout du nœud worker dans le cluster k8s - join

L'opération qui permet au nœud worker de rejoindre le cluster s'appelle le join.

La syntaxe de cette commande est obtenue en lançant la commande suivante sur le control plane avec l'utilisateur root :

```
# kubeadm token create --print-join-command
kubeadm join 10.10.10.30:6443 \
--token ilfbgc.8xco4svm5pnxkfbj \
--discovery-token-ca-cert-hash sha256:73bf45619ae0051d4ff810328d1dadcd18e6a5966c95d3c4ec76275b89a934595
```

Lancement du join sur chacun des workers

Sur chacun des workers, le lancement de la commande join produit le résultat suivant :

```
# kubeadm join 10.10.10.30:6443 \
--token 6pia7c.n6u8pbm7yjl6nnr8 \
--discovery-token-ca-cert-hash sha256:f6d45602ea75c7659dc91f661d19e97e6817e2847e4e5d0047880b871317a145
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system \
get cm kubeadm-config -o yaml'
W0315 16:31:41.445771 6266 configset.go:78] Warning: No kubeproxy.config.k8s.io/v1alpha1 config is loaded.
Continuing without it: configmaps "kube-proxy" is forbidden: User "system:bootstrap:6pia7c"
cannot get resource "configmaps" in API group "" in the namespace "kube-system"
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...
```

This node has joined the cluster:

- * Certificate signing request was sent to apiserver and a response was received.
- * The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

Lancement du join sur chacun des workers

La commande suivante permet de vérifier le résultat du join :

```
$ kubectl get nodes
NAME      STATUS    ROLES      AGE      VERSION
debian-cp  NotReady control-plane  15h      v1.26.2
dnumworker1  NotReady <none>     53s      v1.26.2
```

Stockage

Terminologie du stockage dans k8s

- Le stockage permanent des données s'appuie sur les volumes persistants (PV) (<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>)
- Un PV est un espace de stockage mis à disposition par k8s.
- Il peut être alloué manuellement ou dynamiquement par l'intermédiaire des storage classes (<https://kubernetes.io/docs/concepts/storage/storage-classes/>)
- Les PV sont l'équivalent d'un node dans un cluster.
- Les persistentVolumeClaim (PVC) sont l'équivalent d'un pod.

Déploiement du stockage - Rook Ceph

- Le storage class sur lequel s'appuie l'opérateur PostgreSQL est Ceph
- L'opérateur k8s **Rook Ceph** facilite le déploiement de Ceph
- Le déploiement s'appuie sur le lien
<https://rook.io/docs/rook/v1.9/quickstart.html>
- La version de l'opérateur utilisée est la v1.9
- Elle supporte les versions k8s v1.17+

Prérequis au déploiement de l'opérateur - Rook Ceph

- Le déploiement de l'opérateur scanne l'ensemble des noeuds de stockage pour vérifier la présence de :
 - des devices bruts (sans partitions ou filesystems formattés)
 - des partitions brutes (sans filesystems formattés)
 - les volumes physiques initialisés par LVM

L'exemple ci-dessous indique comment vérifier la disponibilité d'espace pour l'opérateur Rook Ceph :

```
lsblk -f
```

NAME	FSTYPE	LABEL	UUID	MOUNTPOINT
vda				
-vda1	LVM2_member		>eS050t-GkUV-YKTH-WsGq-hNjY-eKNf-3i07IB	
-ubuntu--vg-root	ext4		c2366f76-6e21-4f10-a8f3-6776212e2fe4	/
-ubuntu--vg-swap_1	swap		9492a3dc-ad75-47cd-9596-678e8cf17ff9	[SWAP]
vdb				

Prérequis au déploiement de l'opérateur - Rook Ceph

- Dans l'exemple précédent, si la colonne *FSTYPE* est renseignée, cela indique la présence d'un filesystem
- La partition vdb n'est pas formatée avec un filesystem : elle est donc utilisable par l'opérateur Rook Ceph
- Le paquet **lvm2** est une dépendance importante de Rook Ceph

Sélection des nœuds sur lesquels Ceph sera déployé

L'opérateur Rook Ceph offre la possibilité de sélectionner les nœuds sur lesquels le stockage Ceph est déployé.

Pour cela, il s'appuie sur la notion de label.

Dans le cadre du POC, les 3 nœuds suivants sont sélectionnés pour porter le stockage :

- dnumcephworker1
- dnumcephworker2
- dnumcephworker3

Affectation des labels sur les nœuds de stockage

Depuis le control plane, lancer les commandes suivantes pour marquer les nœuds :

```
$ kubectl label nodes dnumcephworker1 role=storage-node
node/dnumcephworker1 labeled
$ kubectl label nodes dnumcephworker2 role=storage-node
node/dnumcephworker2 labeled
$ kubectl label nodes dnumcephworker3 role=storage-node
node/dnumcephworker3 labeled
```

Affichage du label des nœuds :

```
$ kubectl get nodes --show-labels
NAME      STATUS   ROLES      LABELS
dnumcephworker1  Ready    <none>   kubernetes.io/hostname=dnumcephworker1,kubernetes.io/os=linux,role=storage-node
dnumcephworker2  Ready    <none>   kubernetes.io/hostname=dnumcephworker2,kubernetes.io/os=linux,role=storage-node
dnumcephworker3  Ready    <none>   kubernetes.io/hostname=dnumcephworker3,kubernetes.io/os=linux,role=storage-node
```

Paramétrage pour la répartition du stockage Ceph sur les nœuds labelisés

```
~/rook$ git diff
diff --git a/deploy/examples/cluster.yaml b/deploy/examples/cluster.yaml
index 9bd50ec97..fef3f777f 100644
--- a/deploy/examples/cluster.yaml
+++ b/deploy/examples/cluster.yaml
@@ -154,22 +154,22 @@ spec:
  # To control where various services will be scheduled by kubernetes, use the placement configuration sections below.
  # The example under 'all' would have all services scheduled on kubernetes nodes labeled with 'role=storage-node' and
  # tolerate taints with a key of 'storage-node'.
- # placement:
- #   all:
- #     nodeAffinity:
- #       requiredDuringSchedulingIgnoredDuringExecution:
- #         nodeSelectorTerms:
- #           - matchExpressions:
- #             - key: role
- #               operator: In
- #             values:
- #               - storage-node
- #     podAffinity:
- #     podAntiAffinity:
- #   topologySpreadConstraints:
- #     tolerations:
- #       - key: storage-node
- #         operator: Exists
+ placement:
+   all:
+     nodeAffinity:
+       requiredDuringSchedulingIgnoredDuringExecution:
+         nodeSelectorTerms:
+           - matchExpressions:
+             - key: role
+               operator: In
+             values:
+               - storage-node
+     podAffinity:
+     podAntiAffinity:
+   topologySpreadConstraints:
+   tolerations:
+     - key: storage-node
+       operator: Exists
  # The above placement information can also be specified for mon, osd, and mgr components
  #   mon:
  #     Monitor deployments may contain an anti-affinity rule for avoiding monitor
```

Paramétrage pour la répartition du stockage Ceph sur les nœuds labelisés

La directive *nodeSelectorTerms* permet de sélectionner les noeuds portant la storageclass Ceph

```
+     nodeSelectorTerms:  
...  
+         - storage-node  
+     podAffinity:  
+     podAntiAffinity:
```

Déploiement de l'opérateur Rook Ceph

Comme indiqué dans le lien

<https://rook.io/docs/rook/v1.9/quickstart.html>, l'application des commandes ci-dessous amorce le déploiement de l'opérateur :

```
$ git clone --single-branch --branch v1.9.2 https://github.com/rook/rook.git  
cd rook/deploy/examples  
kubectl create -f crds.yaml -f common.yaml -f operator.yaml  
kubectl create -f cluster.yaml
```

- Une fois le cluster opérationnel, il devient possible de créer :
 - stockage bloc
 - stockage objet
 - stockage fichier

Vérification de l'opérateur Rook Ceph

```
# verify the rook-ceph-operator is in the `Running` state before proceeding
kubectl -n rook-ceph get pod
NAME                               READY   STATUS    RESTARTS   AGE
csi-cephfsplugin-9nbts            2/2     Running   1 (63d ago)  69d
csi-cephfsplugin-bpxlw            2/2     Running   0          69d
csi-cephfsplugin-jd5x8            2/2     Running   0          69d
csi-cephfsplugin-mddkf            2/2     Running   0          69d
csi-cephfsplugin-nrmfz            2/2     Running   0          69d
csi-cephfsplugin-provisioner-84cc595b78-9mml4  5/5     Running   6523 (28d ago)  69d
csi-cephfsplugin-provisioner-84cc595b78-9twnq  5/5     Running   3908 (38d ago)  69d
csi-rbdplugin-92z1q              2/2     Running   0          69d
csi-rbdplugin-c95w7              2/2     Running   0          69d
csi-rbdplugin-pk57s              2/2     Running   1 (63d ago)  63d
csi-rbdplugin-provisioner-6f6b6b8cd6-4c8jd  5/5     Terminating  2919 (29d ago)  69d
csi-rbdplugin-provisioner-6f6b6b8cd6-d4t56  0/5     Pending    0          4d10h
csi-rbdplugin-provisioner-6f6b6b8cd6-gw6bm  1/5     CreateContainerError  4465  69d
csi-rbdplugin-srtfz              2/2     Running   0          69d
csi-rbdplugin-v6gqm              2/2     Running   0          69d
rook-ceph-crashcollector-dnumcephworker1-7845bb8ff-vs9fx  1/1     Running   0          68d
rook-ceph-crashcollector-dnumcephworker2-75cdf95dcn-ljkqd  0/1     Pending    0          4d10h
rook-ceph-crashcollector-dnumcephworker2-75cdf95dcn-n5xsz  1/1     Terminating  0          69d
rook-ceph-crashcollector-dnumcephworker3-6fdbb6cd9-x45w5  1/1     Running   2          68d
rook-ceph-mgr-a-c5db58dff-fpp7z  2/3     CrashLoopBackOff  146 (28d ago)  30d
rook-ceph-mgr-a-c5db58dff-hvsp9  2/3     Terminating  3115 (30d ago)  69d
rook-ceph-mgr-b-7bbfd88c8b-jdg4p  0/3     Pending    0          4d10h
rook-ceph-mgr-b-7bbfd88c8b-wh4ww  2/3     Terminating  2283 (28d ago)  58d
rook-ceph-mon-a-75cf9ccddc-b2jgc  2/2     Running   1500 (31d ago)  69d
rook-ceph-mon-c-64dcba86c-wz8sg  2/2     Running   1808 (28d ago)  69d
rook-ceph-operator-cf4f7df4-6tm6p  1/1     Running   0          68d
rook-ceph-osd-0-57d9b8db4d-d6dhr  1/2     Terminating  731 (28d ago)  68d
rook-ceph-osd-0-57d9b8db4d-vmtjp  0/2     Pending    0          4d10h
rook-ceph-osd-1-7469877fd-6n2mh  1/2     Running   716 (30d ago)  68d
rook-ceph-osd-2-5cc486467c-lhm47  1/2     Running   1172 (28d ago)  68d
rook-ceph-osd-prepare-dnumcephworker1-rnk78  0/1     Completed  0          57d
rook-ceph-osd-prepare-dnumcephworker3-42rxv  0/1     Completed  0          57d
rook-ceph-tools-7c4b8bb9b5-8tf8r  0/1     Pending    0          4d10h
rook-ceph-tools-7c4b8bb9b5-pxk67  1/1     Terminating  0          68d
```

Le contrôleur d'admission (Admission Controller) - Rook Ceph

- Il est recommandé de déployer le contrôleur d'admission : il permet de vérifier que Rook est correctement paramétré grâce aux réglages des Customer Resources (CR)
- L'Admission Controller intercepte les requêtes à destination de l'API k8s avant l'objet persistant après les phases d'authentification et d'autorisation
- Pour installer l'Admission Controller, lancer les requêtes suivantes :
`kubectl apply -f https://github.com/jetstack/cert-manager/releases/download/v1.7.1/cert-manager.yaml`

Affichage des storage class déployés

Le storageclass déployé a pour nom **rook-ceph-block**.

linagora@debian-cp:~\$ kubectl get storageclass						
NAME	PROVISIONER	RECLAIMPOLICY	VOLUMEBINDINGMODE	ALLOWVOLUMEEXPANSION	AGE	
local-storage	kubernetes.io/no-provisioner	Delete	WaitForFirstConsumer	false	12d	
rook-ceph-block	rook-ceph.rbd.csi.ceph.com	Delete	Immediate	true	5d23h	

Déploiement PostgreSQL

Sélection des nœuds PostgreSQL

De manière similaire à l'opérateur Rook Ceph, il est possible de sélectionner les nœuds portant le pod PostgreSQL en se basant sur les labels Kubernetes.

Marquage des nœuds PostgreSQL

Les commandes ci-dessous marquent les nœuds destinés à porter les pods PostgreSQL :

```
$ kubectl label nodes dnumworker1 postgres-operator=enabled
node/dnumworker1 labeled
$ kubectl label nodes dnumworker2 postgres-operator=enabled
node/dnumworker2 labeled
$ kubectl get nodes --show-labels
NAME           STATUS   ROLES      LABELS
dnumworker1    Ready    <none>     kubernetes.io/hostname=dnumworker1,kubernetes.io/os=linux,postgres-operator=enabled
dnumworker2    Ready    <none>     kubernetes.io/hostname=dnumworker2,kubernetes.io/os=linux,postgres-operator=enabled
```

Répartitions des pods PostgreSQL sur les nœuds worker et choix du storageClass

```
$ git diff
diff --git a/manifests/complete-postgres-manifest.yaml b/manifests/complete-postgres-manifest.yaml
index 8d197a75..56b32c34 100644
--- a/manifests/complete-postgres-manifest.yaml
+++ b/manifests/complete-postgres-manifest.yaml
@@ -57,7 +57,7 @@ spec:

volume:
  size: 1Gi
-#   storageClass: my-sc
+  storageClass: rook-ceph-block
 #   iops: 1000 # for EBS gp3
 #   throughput: 250 # in MB/s for EBS gp3
 #   selector:
@@ -203,14 +203,14 @@ spec:

# Add node affinity support by allowing postgres pods to schedule only on nodes that
# have label: "postgres-operator:enabled" set.

...
+  nodeAffinity:
+    requiredDuringSchedulingIgnoredDuringExecution:
+      nodeSelectorTerms:
+        - matchExpressions:
+          - key: postgres-operator
+            operator: In
+            values:
+              - enabled

# Enables change data capture streams for defined database tables
# streams:
```

Déploiement de l'opérateur PostgreSQL de Zalando

Le storage class est maintenant déployé.

Il devient possible d'appliquer l'opérateur PostgreSQL.

Le lien suivant² décrit les commandes à appliquer.

2.

<https://github.com/zalando/postgres-operator/blob/master/docs/quickstart.md#deployment-options>

Clonage du dépôt de l'opérateur

```
git clone https://github.com/zalando/postgres-operator.git  
cd postgres-operator
```

Application des différents manifestes

```
kubectl create -f manifests/configmap.yaml # configuration
kubectl create -f manifests/operator-service-account-rbac.yaml # identity and permissions
kubectl create -f manifests/postgres-operator.yaml # deployment
kubectl create -f manifests/api-service.yaml # operator API to be used by UI
```

Pour information, il existe également des chart Helm pour facilier le déploiement.

Accès à l'interface web

Pour activer l'accès à l'interface web de l'opérateur PostgreSQL, veuillez entrer la commande suivante sur le nœud control plane :

```
$ kubectl port-forward svc/postgres-operator-ui 8081:80
Forwarding from 127.0.0.1:8081 -> 8081
Forwarding from [::1]:8081 -> 8081
```

Elle redirige le flux TCP du port 80 du control plane vers le port TCP 8081 du service postgres-operator-ui

Accès à l'interface web

Pour accéder à l'interface web de l'opérateur PostgreSQL depuis le PC de l'utilisateur, il est possible de passer par une redirection SSH :

```
ssh -L 9090:10.106.57.137:80 dgfip-k8s
```

Lancer le navigateur pour accéder à l'URL <http://localhost:9090/#new>

Interface web de l'opérateur PostgreSQL

The screenshot shows the PostgreSQL Operator UI interface. At the top, there is a navigation bar with links for PostgreSQL Operator UI, PostgreSQL clusters, Backups, Status, New cluster, and Documentation. The title "New PostgreSQL cluster" is displayed prominently. On the left, there is a "Cluster YAML definition" section containing the YAML configuration for the cluster. On the right, there is a "New cluster configuration" form with various input fields and checkboxes.

Cluster YAML definition

```
kind: "postgresql"
apiVersion: "acid.zalan.do/v1"

metadata:
  name: ""
  namespace: "default"
  labels:
    team: acid

spec:
  teamId: "acid"
  postgresql:
    version: "15"
    numberofInstances: 1
  volume:
    size: "10Gi"

allowedSourceRanges:
  # IP ranges to access your cluster go here

resources:
  requests:
    cpu: 100m
    memory: 100Mi
  limits:
    cpu: 500m
    memory: 500Mi
```

New cluster configuration

New cluster configuration	
Name	new-cluster (can be 53 characters long)
Owning team	acid
PostgreSQL version	15
DNS name:	.default
Number of instances	1
Enable load balancer	<input type="checkbox"/> Master
Enable connection pooler	<input type="checkbox"/> Master
Enable connection pooler load balancer	<input type="checkbox"/> Master
Volume size	10 Gi
storageClass	
lops	

Below the "storageClass" field, there is a note: "Specify lops and Throughput only if you need more than the default 3000 lops and 125Mb/s EBS provides."

Fonctionnalités proposées par l'interface web de l'opérateur PostgreSQL

L'UI permet de :

- choisir la version PostgreSQL (jusqu'à la version 15 actuellement)
- le nombre d'instances
- activation du load-balancer
- activation du pool de connexions à la base
- activation du load-balancer pour le pool de connexions à la base
- taille du volume persistent alloué à la base de données
- choix du storageClass
- performances IO
- choix des ressources (demandées et limites) CPU et RAM allouées

Utilisation de la commande en ligne pour la création d'un cluster PostgreSQL

- Les fonctionnalités proposées par l'UI sont également disponibles par la commande en ligne.
- Le manifeste *manifests/complete-postgresql-manifest.yaml* permet de préciser l'ensemble des paramètres proposés par l'UI.
- Pour appliquer ce manifeste *manifests/complete-postgresql-manifest.yaml*, la commande suivante est lancée sur le nœud :

```
kubectl create -f manifests/complete-postgresql-manifest.yaml
```

Vérification de l'état du cluster PostgreSQL

```
$ kubectl get pods -l application=spilo -L spilo-role
NAME           READY   STATUS    RESTARTS   AGE   SPILO-ROLE
acid-test-cluster-0  1/1    Running   0          6m49s  master
acid-test-cluster-1  1/1    Running   0          6m12s  replica
$ kubectl get postgresql
NAME      TEAM  VERSION  PODS  VOLUME  CPU-REQUEST  MEMORY-REQUEST  AGE  STATUS
acid-test-cluster  acid   15       2     1Gi     10m          100Mi        68d  Running
```

Stockage S3 - Minio

Déploiement de krew

L'opérateur Minio s'appuie sur le gestionnaire de paquets **krew**.

```
(  
    set -x; cd "$(mktemp -d)" &  
    OS="$(uname | tr '[:upper:]' '[:lower:]')" &  
    ARCH="$(uname -m | sed -e 's/x86_64/amd64/' -e 's/^(arm\)(64)\?.*$/\1\2/' -e 's/aarch64$/arm64/')" &  
    KREW="krew-${OS}_${ARCH}" &  
    curl -fsSLO "https://github.com/kubernetes-sigs/krew/releases/latest/download/${KREW}.tar.gz" &  
    tar zxvf "${KREW}.tar.gz" &  
    ./"${KREW}" install krew  
)
```

Déploiement de krew

Ajout du répertoire des binaires du paquet krew dans **.bashrc** ou **.zshrc** :

```
export PATH="${KREW_ROOT:-$HOME/.krew}/bin:$PATH"
```

Redémarrer le shell.

Pour vérifier le déploiement correct de krew, lancer la commande suivante :

```
kubectl krew
```

Déploiement de l'opérateur Minio

Le déploiement de Minio s'appuie sur l'opérateur Minio.

Son déploiement est décrit dans le lien suivant

<https://operator.min.io/#architecture>.

```
kubectl krew update  
kubectl krew install minio
```

Vérification de l'état de l'opérateur Minio

```
$ kubectl get pods -n minio-operator
NAME                  READY   STATUS    RESTARTS   AGE
console-56f9795d5c-59fsx   1/1    Running   1 (33d ago)   82d
minio-operator-7cd6784f59-5c52w   1/1    Running   5 (23h ago)   17d
minio-operator-7cd6784f59-m7h8x   1/1    Running   2320 (3d16h ago) 82d
```

Accès à la console Minio

La commande suivante ouvre un accès de type proxy à la console Minio :

```
$ kubectl minio proxy -n minio-operator
Starting port forward of the Console UI.

To connect open a browser and go to http://localhost:9090

Current JWT to login: *****

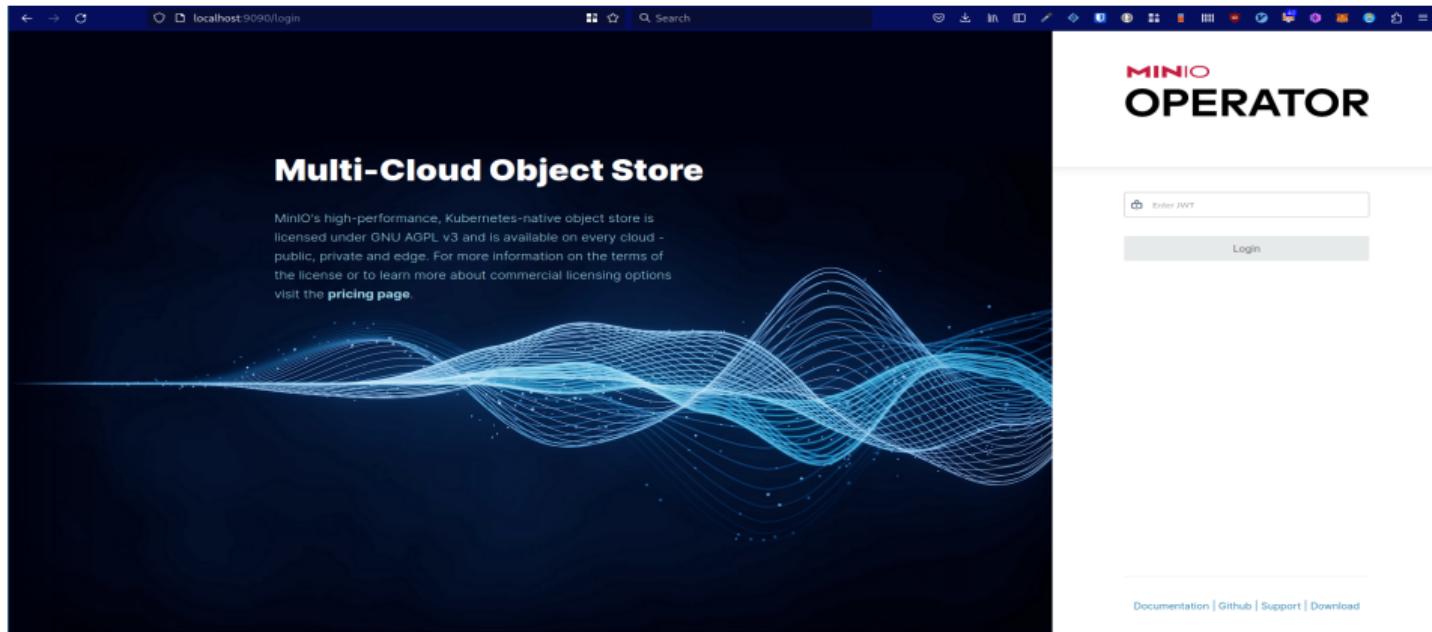
Forwarding from 0.0.0.0:9090 -> 9090
Handling connection for 9090
```

Depuis le terminal de l'utilisateur, lancer la commande suivante :

```
$ ssh -L 9090:localhost:9090 dgfip-k8s
```

Accès à la console Minio

Dans le champ *Enter JWT*, renseigner la valeur du token JWT renvoyé par la commande précédente :



Dashboard Minio

Le tableau de bord de Minio ressemble à ceci :

The screenshot shows a web browser window with the URL `localhost:9090/tenants`. The title bar includes standard icons for back, forward, search, and refresh. The main header has the "MINIO OPERATOR" logo. On the left, a sidebar menu lists "Tenants" (selected), "License", "Register", and "Documentation". The main content area is titled "Tenants" and contains a sub-section titled "Tenants". It explains that a tenant is a logical structure for a Minio deployment, noting size and configuration differences from other tenants. A call-to-action button "Create a Tenant" is visible. At the bottom left of the content area, there is a "Sign Out" link.

Création d'un tenant - Setup - Minio

The screenshot shows the Minio Operator web interface for creating a new tenant. The left sidebar has links for License, Register, and Documentation. The main area is titled "Tenants" and shows a "Setup" configuration page.

Name
How would you like to name this new tenant?
Name* [Input field]

Configure

Images
Namespace* [Input field]

Pod Placement
Storage Class [Input field]

Capacity
Please select the desired capacity

Identity Provider

Security
Number of Servers* [Input field] 4

Encryption
Drives per Server* [Input field] 4

Audit Log
Total Size* [Input field] 1024 Gi

Monitoring

Erasure Code Parity EC-4 (Default)

Please select the desired parity. This setting will change the max usable capacity in the cluster

Resources
You may specify the amount of CPU and Memory that MinIO servers should reserve on each node.

CPU Request [Input field] 1

Memory Request [Input field] 2 Gi

Specify Limit OFF ON

Resource Allocation

Number of Servers	4
Drives per Server	4
Drive Capacity	64.0 GB
Total Volumes	16
Memory per Node	2 Gi
CPU Selection	1

Erasure Code Configuration

EC Parity	EC-4
Raw Capacity	1.0 TiB
Usable Capacity	768.0 GiB
Server Failures Tolerated	1

Création d'un tenant - Configure - Minio

The screenshot shows the Minio Operator web interface. On the left is a sidebar with the following options:

- Minio OPERATOR
- ← Tenants
- License
- Register
- Documentation

The main content area has a header "Tenants" with a back arrow. Below it is a table with two columns:

Setup	Configure
	Configure Basic configurations for tenant management
	Services Whether the tenant's services should request an external IP via LoadBalancer service type.
Images	Expose MinIO Service <input checked="" type="checkbox"/> ON
Pod Placement	Expose Console Service <input checked="" type="checkbox"/> ON
Identity Provider	Set Custom Domains <input checked="" type="checkbox"/> ON
Security	Security Context <input checked="" type="checkbox"/> ON
Encryption	
Audit Log	
Monitoring	

Below the table, there is a section for **Additional Environment Variables** with a note: "Define additional environment variables to be used by your MinIO pods". It includes a table for adding key-value pairs:

Key	Value	+	-

At the bottom right are "Cancel" and "Create" buttons.

Marquage des nœuds Minio

Depuis le control plane, lancer les commandes suivantes pour marquer les nœuds :

```
$ kubectl label nodes dnumminioworker1 role=s3-node  
node/dnumminioworker1 labeled  
$ kubectl label nodes dnumminioworker2 role=s3-node  
node/dnumminioworker2 labeled  
$ kubectl label nodes dnumminioworker3 role=s3-node  
node/dnumminioworker3 labeled  
$ kubectl label nodes dnumminioworker4 role=s3-node  
node/dnumminioworker4 labeled
```

Marquage des nœuds Minio

Affichage du label des nœuds :

```
$ kubectl get nodes --show-labels
NAME           STATUS    ROLES      AGE     VERSION   LABELS
...
dnumminioworker1  Ready    <none>    6h2m    v1.27.2   ... ,role=s3-node
dnumminioworker2  Ready    <none>    142m    v1.27.2   ... ,role=s3-node
dnumminioworker3  Ready    <none>    109m    v1.27.2   ... ,role=s3-node
dnumminioworker4  Ready    <none>    83m     v1.27.2   ... ,role=s3-node
...
...
```

Création d'un tenant - Pod placement - Minio

The screenshot shows the MinIO Operator web interface at `localhost:9090/tenants/add`. The left sidebar has a dark theme with the following navigation items:

- Tenants (selected)
- License
- Register
- Documentation

The main content area is titled "Tenants" and shows the "Pod Placement" configuration for a new tenant. A red oval highlights the "Labels" section where a role label "s3-node" is assigned to a node.

Pod Placement
Configure how pods will be assigned to nodes

Type
MinIO supports multiple configurations for Pod Affinity

None
 Default (Pod Anti-Affinity)
 Node Selector

With Pod Anti-Affinity OFF ON

Labels

role	s3-node	+	-
------	---------	---	---

Tolerations

Toleration 1

if	Toleration Key	is	Equal	to	Toleration Value
then	NoSchedule	after	0	seconds	+

Cancel Create

Création d'un tenant - Choix d'un fournisseur d'identité - Minio

The screenshot shows the Minio Operator web interface. On the left is a sidebar with the following navigation items:

- MINIO OPERATOR
- Tenants
- License
- Register
- Documentation

The main content area has a breadcrumb navigation bar: « Tenants ». The central panel is titled 'Identity Provider' with the sub-instruction: 'Access to the tenant can be controlled via an external Identity Manager.' Below this, there is a 'Protocol' section with three options: 'Built-in' (selected), 'OpenID', and 'LDAP / Active Directory'. Under 'Add additional users', two user IDs are listed: '7IE7nFpNzV2TlkHq' and 'cMz4n5CEpoWeU9MPmkTcazYix1GrFQ2R'. At the bottom right of this section are 'Cancel' and 'Create' buttons.

Création d'un tenant - Sécurité - Minio

The screenshot shows the Minio Operator web interface for creating a new tenant. The left sidebar has a 'Tenants' tab selected. The main content area is titled 'Tenants' and shows a 'Security' configuration section.

Security

- TLS**: Securing all the traffic using TLS. This is required for Encryption Configuration. Status: ON
- AutoCert**: The internode certificates will be generated and managed by Minio Operator. Status: ON
- Custom Certificates**: Certificates used to terminated TLS at Minio. Status: OFF

Buttons: Cancel, Create

Création d'un tenant - Chiffrement - Minio

The screenshot shows the Minio Operator web interface. The left sidebar has a dark blue background with the Minio logo at the top, followed by sections: **OPERATOR**, **Tenants** (selected), **License**, **Register**, and **Documentation**. Below these are two large circular icons: one with a gear and another with a document. At the bottom is a "Sign Out" button.

The main content area has a light gray header with the text "**← Tenants**". The main panel title is **Encryption**. To the right of the title is a toggle switch labeled "Disabled" (grayed out) and "Enabled". A descriptive text block below the title states: "MinIO Server-Side Encryption (SSE) protects objects as part of write operations, allowing clients to take advantage of server processing power to secure objects at the storage layer (encryption-at-rest). SSE also provides key functionality to regulatory and compliance requirements around secure locking and erasure." At the bottom right of the panel are "Cancel" and "Create" buttons.

A vertical sidebar on the left lists several configuration options: **Setup**, **Configure**, **Images**, **Pod Placement**, **Identity Provider**, **Security**, **Encryption** (which is selected and highlighted in light blue), **Audit Log**, and **Monitoring**.

Création d'un tenant - Log d'audit - Minio

The screenshot shows the Minio Operator web interface. On the left, there's a sidebar with navigation links: 'Tenants' (selected), 'License', 'Register', and 'Documentation'. Below the sidebar, there's a 'Sign Out' button.

The main content area has a title 'Tenants' with a back arrow. It displays a table with a single row for the 'Audit Log' tenant. The table columns are 'Setup', 'Configure', 'Images', 'Pod Placement', 'Identity Provider', 'Security', 'Encryption', 'Audit Log' (which is selected), and 'Monitoring'. The 'Audit Log' row contains several configuration fields:

- Audit Log**: A toggle switch labeled 'Enabled' is turned on.
- Log Search Storage Class**: Set to 'Default'.
- Storage Size***: Set to '\$'.
- SecurityContext for LogSearch**:
 - Run As User***: Set to '1000'.
 - Run As Group***: Set to '1000'.
 - FsGroup***: Set to '1000'.
 - FsGroupChangePolicy**: Set to 'Always'.
 - Do not run as Root**: A toggle switch is turned on.
- SecurityContext for PostgreSQL**:
 - Run As User***: Set to '999'.
 - Run As Group***: Set to '999'.
 - FsGroup***: Set to '999'.
 - FsGroupChangePolicy**: Set to 'Always'.
 - Do not run as Root**: A toggle switch is turned on.

At the bottom right of the configuration area are 'Cancel' and 'Create' buttons.

Création d'un tenant - Supervision - Minio

The screenshot shows the Minio Operator web interface for creating a new tenant. The left sidebar has links for Tenants, License, Register, and Documentation. The main area shows a list of setup steps: Setup, Configure, Images, Pod Placement, Identity Provider, Security, Encryption, Audit Log, and Monitoring. The Monitoring step is currently selected. The monitoring configuration includes:

- Monitoring**: A small Prometheus will be deployed to keep metrics about the tenant. A toggle switch is set to **Enabled**.
- Storage Class**: Default
- Storage Size***: 5 (with units G and T)
- SecurityContext**
 - Run As User***: 1000
 - Run As Group***: 1000
 - FsGroup***: 1000
 - FsGroupChangePolicy**: Always
- Do not run as Root**: A toggle switch is set to **ON**.

At the bottom right are **Cancel** and **Create** buttons.

Image Docker utilisée par l'opérateur PostgreSQL

L'image déployée par l'opérateur PostgreSQL de Zalando s'appuie sur Spilo
<https://github.com/zalando/spilo>

Cette information se trouve dans le script
manifests/complete-postgres-manifest.yaml :

```
dockerImage: ghcr.io/zalando/spilo-15:3.0-p1
```

Sécurité - Support d'OpenShift

Utilisation de container en mode rootless dans OpenShift

- L'URL suivante décrit la problématique d'utilisation de container dans un environnement rootless³
- Cette page fournit également un lien intéressant sur la sécurisation d'un cluster Kubernetes en se basant sur les **Pod Security Policies**.
- Le lien est le suivant : <https://docs.bitnami.com/tutorials/secure-kubernetes-cluster-psp/>.

3.

<https://docs.bitnami.com/tutorials/running-non-root-containers-on-openshift>

Pod Security Policies

Comme indiqué dans <https://kubernetes.io/docs/concepts/security/pod-security-policy/>:

//kubernetes.io/docs/concepts/security/pod-security-policy/,
les PSP sont dépréciés.

Ils sont maintenant remplacés par Pod Security Admission <https://kubernetes.io/docs/concepts/security/pod-security-admission/>.

La norme *Pod Security Admission* définit la notion de *Security Context* décrite
dans le lien suivant : <https://kubernetes.io/docs/tasks/configure-pod-container/security-context/> C'est cette notion qui va
permet d'approcher au plus près les conditions de run d'un cluster Openshift

Opérateur - Security Context

Par défaut, l'opérateur PostgreSQL de Zalando applique les mesures de sécurité suivantes.

Extrait de manifests/postgres-operator.yaml :

```
securityContext:  
  runAsUser: 1000  
  runAsNonRoot: true  
  readOnlyRootFilesystem: true  
  allowPrivilegeEscalation: false
```

- Il ne tourne pas avec un identifiant privilégié ni le compte root
- Il s'appuie sur des filesystems en lecture seule

Pods PostgreSQL - Security Context

Il est possible d'affecter :

- un utilisateur non privilégié
- un groupe non privilégié au pod.
- un groupe de filesystem défini

Extrait de manifests/complete-postgres-manifest.yaml :

```
spiloRunAsUser: 101
spiloRunAsGroup: 103
spiloFSGroup: 103
```

Déploiement continu

Installation d'ArgoCD

Le lien suivant décrit l'installation d'ArgoCD :

<https://argo-cd.readthedocs.io/en/stable/#quick-start>

```
kubectl create namespace argocd  
kubectl apply -n argocd -f https://raw.githubusercontent.com/argoproj/argo-cd/stable/manifests/install.yaml
```

Installation de la CLI ArgoCD

Le lien suivant⁴ décrit l'installation de la CLI ArgoCD.

Les commandes d'installation sont :

```
curl -sSL -o argocd-linux-amd64 https://github.com/argoproj/argo-cd/releases/latest/download/argocd-linux-amd64  
sudo install -m 555 argocd-linux-amd64 /usr/local/bin/argocd  
rm argocd-linux-amd64
```

4.

https://argo-cd.readthedocs.io/en/stable/cli_installation/#download-with-curl

Activation de l'accès au serveur d'API d'ArgoCD

Le lien suivant⁵ décrit les différentes méthodes d'accès au serveur d'API d'ArgoCD :

Les commandes d'installation sont :

```
$ kubectl patch svc argocd-server -n argocd -p '{"spec": {"type": "LoadBalancer"}}'  
service/argocd-server patched
```

5.

https://argo-cd.readthedocs.io/en/stable/getting_started/#3-access-the-argo-cd-api-server

Build de l'image Docker Spilo

Le Dockerfile définissant l'image Spilo est disponible à l'URL suivante :

<https://github.com/zalando/spilo>

La méthode de génération de l'image est décrite dans : <https://github.com/zalando/spilo#how-to-build-this-docker-image>

```
$:~/spilo/postgres-appliance$ docker build --tag dnum-test .
[+] Building 9762.1s (33/33) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 3.01kB
=> [internal] load .dockerignore
=> => transferring context: 2B
...
=> => exporting layers
=> => writing image sha256:52afd69aff9414333220ec408283a8ff20e2162e703c6ab5af5090d9d62e4e0
=> => naming to docker.io/library/dnum-test
```

Le build dure un peu moins de **2h43min.**

Login avec la CLI d'ArgoCD

La commande suivante permet de récupérer le mot de passe initial de l'admin d'ArgoCD :

```
linagora@debian-cp:~$ argocd admin initial-password -n argocd  
*****
```

```
This password must be only used for first time login. We strongly recommend you  
update the password using `argocd account update-password`.
```

Login avec la CLI d'ArgoCD

```
$ kubectl get svc -n argocd
NAME                      TYPE        CLUSTER-IP      EXTERNAL-IP    PORT(S)          AGE
argocd-applicationset-controller   ClusterIP  10.109.27.242  <none>        7000/TCP,8080/TCP  6d16h
argocd-dex-server                ClusterIP  10.100.63.110  <none>        5556/TCP,5557/TCP,5558/TCP  6d16h
argocd-metrics                  ClusterIP  10.102.19.237  <none>        8082/TCP         6d16h
argocd-notifications-controller-metrics ClusterIP  10.102.70.229  <none>        9001/TCP         6d16h
argocd-redis                     ClusterIP  10.104.253.209  <none>        6379/TCP         6d16h
argocd-repo-server               ClusterIP  10.111.28.234  <none>        8081/TCP,8084/TCP  6d16h
argocd-server                    LoadBalancer  10.109.47.144  <pending>    80:30235/TCP,443:30885/TCP  6d16h
argocd-server-metrics             ClusterIP  10.98.167.31   <none>        8083/TCP         6d16h
linagora@debian-cp:~$ argocd login 10.109.47.144
WARNING: server certificate had error: x509: cannot validate certificate for 10.109.47.144 because it doesn't contain
any IP SANs. Proceed insecurely (y/n)? y
Username: admin
Password:
'admin:login' logged in successfully
Context '10.109.47.144' updated
```

Ajout de l'application postgres-operator dans ArgoCD

- ArgoCD se met à l'écoute des changements d'un projet de déploiement dans un dépôt.
- L'étape suivante "abonne" ArgoCD au dépôt Github sur lequel est enregistré le déploiement de l'opérateur postgres
- La commande qui lie ArgoCD au dépôt de l'opérateur est la suivante :

```
argocd app create postgres-operator \
--repo https://github.com/simonelbaz/postgres-operator.git \
--path manifests \
--dest-server https://kubernetes.default.svc \
--revision poc-argocd \
--dest-namespace default
```

Ajout de l'application postgres-operator dans ArgoCD

- Cette commande est décrite en détail dans la documentation officielle⁶
- Pour information, le dépôt ci-dessus est un fork du dépôt officiel
- L'ensemble des modifications réalisées pour le POC sont tracées dans la branche *poc-argocd*

6.

[https://argo-cd.readthedocs.io/en/stable/user-guide/commands/
argocd_app_create/](https://argo-cd.readthedocs.io/en/stable/user-guide/commands/argocd_app_create/)

Ajout de l'application postgres-operator dans ArgoCD

- Pour vérifier que l'application a correctement été ajoutée dans ArgoCD :

```
$ argocd app list
NAME                CLUSTER          NAMESPACE PROJECT STATUS HEALTH SYNCPOLICY
argocd/postgres-operator https://kubernetes.default.svc default default Unknown Healthy <none>
```

- Selon le lien suivant https://argo-cd.readthedocs.io/en/stable/user-guide/tool_detection/, le script kustomization.yaml est automatiquement détecté par ArgoCD

- ArgoCD est capable de s'interfacer avec Kustomize
<https://kustomize.io/>.
- Zalando met à disposition un script Yaml `manifests/kustomization.yaml`

```
$ cat kustomization.yaml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
resources:
- configmap.yaml
- operator-service-account-rbac.yaml
- postgres-operator.yaml
- api-service.yaml
```

- Ce dernier facilite le déploiement de l'opérateur dans l'environnement k8s
- Une page synthétisant l'utilisation de Kustomize est disponible à l'URL suivante : <https://kubectl.docs.kubernetes.io/guides/introduction/kustomize/>

- Le lien suivant décrit l'intégration d'un outil de CI/CD et ArgoCD.
https://argo-cd.readthedocs.io/en/stable/user-guide/ci_automation/
- Globalement la mise à jour d'un cluster k8s par l'intermédiaire d'ArgoCD se déroule en 2 phases :
 - récupération du dépôt git
 - la phase de patch
 - la phase de sync

Phase de patch & ArgoCD

- La 1^{re} étape de la phase de patch est de récupérer le dépôt du projet :
`git clone https://github.com/simonelbaz/postgres-operator.git
cd postgres-operator`
- Il devient possible de patcher avec la commande **kustomize** ou la commande **kubectl**
 - Patch avec kustomize (nécessite l'installation de kustomize) :

Phase de patch avec kustomize

- L'URL suivante décrit l'installation de kustomize : <https://kubectl.docs.kubernetes.io/installation/kustomize/binaries/>
- Pour installer kustomize, merci de lancer la commande suivante :
`curl -s "https://raw.githubusercontent.com/kubernetes-sigs/kustomize/master/hack/install_kustomize.sh" | bash`
- kustomize peut servir à l'édition de kustomization.yaml.
- Dans le cadre du POC, cela n'a pas été nécessaire

Phase de patch avec kubectl

- On liste dans un 1^{er} temps l'ensemble des ConfigMaps :

```
$ kubectl get configmaps  
NAME          DATA   AGE  
kube-root-ca.crt  1      98d  
postgres-operator 59     97d
```

- Pour lister la ConfigMap qui définit l'image **spilo** déployée :

```
$ kubectl get configmap postgres-operator -o yaml  
apiVersion: v1  
data:  
  api_port: "8080"  
  ...  
  docker_image: ghcr.io/zalando/spilo-15:2.1-p9  
  ...  
kind: ConfigMap  
metadata:  
  creationTimestamp: "2023-03-16T22:20:28Z"  
  name: postgres-operator  
  namespace: default  
  ...
```

Phase de patch avec kubectl

La version de l'image spilo est modifiée :

```
kubectl patch --local -f configmap.yaml \
-p '{"data":{"docker_image":"ghcr.io/zalando/spilo-15:3.0-p1"}}' \
-o yaml > configmap.yaml.new
mv configmap.yaml.new configmap.yaml
```

Phase de sync & ArgoCD

- Une fois la modification réalisée, il est nécessaire de la pusher vers le dépôt sur lequel se synchronise ArgoCD

```
git add . -m "Mise à jour de l'image"  
git push
```

- Lancer ensuite les commandes suivantes dans le pipeline de la CI :

```
export ARGOCD_SERVER=argocd.mycompany.com  
export ARGOCD_AUTH_TOKEN=<JWT token generated from project>  
curl -sSL -o /usr/local/bin/argocd https://$ARGOCD_SERVER/download/argocd-linux-amd64  
argocd app sync nom application  
argocd app wait nom application
```

ArgoCD - Etat du cluster avant le lancement du sync

```
$ git log | head -1
commit 90ad8c7aed92c6430bcb36ab23228a2d57c2715d
$ argocd app get argocd/postgres-operator
Name:          argocd/postgres-operator
Project:       default
Server:        https://kubernetes.default.svc
Namespace:     default
URL:          https://10.109.47.144/applications/postgres-operator
Repo:          https://github.com/simonelbaz/postgres-operator.git
Target:        poc-argocd
Path:          manifests
SyncWindow:    Sync Allowed
Sync Policy:   <none>
Sync Status:   OutOfSync from poc-argocd (90ad8c7)
Health Status: Healthy
```

GROUP	KIND	NAMESPACE	NAME	STATUS	HEALTH	HOOK	MESSAGE
acid.zalan.do	ConfigMap	default	postgres-operator	OutOfSync	Healthy		
	Service	default	postgres-operator	OutOfSync			
	ServiceAccount	default	postgres-operator	OutOfSync			
apps	postgresql	default	acid-test-cluster	OutOfSync			
rbac.authorization.k8s.io	Deployment	default	postgres-operator	OutOfSync	Healthy		
rbac.authorization.k8s.io	ClusterRole		postgres-operator	OutOfSync			
rbac.authorization.k8s.io	ClusterRole		postgres-pod	OutOfSync			
rbac.authorization.k8s.io	ClusterRoleBinding		postgres-operator	OutOfSync			

ArgoCD - Etat du cluster avant le lancement du sync

- L'information renvoyée par la commande `argocd get` est cohérente.
- Il y a bien un décalage entre l'état du cluster et le dernier commit de la branche poc-argocd
- Vérification de la version de l'image spilo déployée :

```
$ kubectl describe pod acid-test-cluster-1 | grep spilo
Labels:           application=spilo
Image:          ghcr.io/zalando/spilo-15:2.1-p9
KUBERNETES_ROLE_LABEL:    spilo-role
KUBERNETES_LABELS:      "application":"spilo"
```

ArgoCD - Lancement de la synchronisation

```
$ argocd app sync postgres-operator
TIMESTAMP          GROUP           KIND      NAMESPACE   NAME    STATUS  HEALTH   HOOK   MESSAGE
2023-06-23T16:49:22+02:00 rbac.authorization.k8s.io ClusterRoleBinding   default  postgres-operator  OutOfSync
2023-06-23T16:49:22+02:00          ConfigMap        default  postgres-operator  OutOfSync
2023-06-23T16:49:22+02:00          Service         default  postgres-operator  OutOfSync Healthy
2023-06-23T16:49:22+02:00          ServiceAccount   default  postgres-operator  OutOfSync
2023-06-23T16:49:22+02:00 acid.zalan.do    postgresql     default  acid-test-cluster  OutOfSync
2023-06-23T16:49:22+02:00 apps            Deployment    default  postgres-operator  OutOfSync Healthy
2023-06-23T16:49:22+02:00 rbac.authorization.k8s.io ClusterRole   default  postgres-operator  OutOfSync
2023-06-23T16:49:22+02:00          ClusterRole     default  postgres-pod   OutOfSync
2023-06-23T16:49:52+02:00          ServiceAccount   default  postgres-operator  Synced
2023-06-23T16:49:52+02:00          ConfigMap        default  postgres-operator  Synced
2023-06-23T16:49:53+02:00 rbac.authorization.k8s.io ClusterRole   default  postgres-operator  Synced
2023-06-23T16:49:53+02:00          ClusterRole     default  postgres-pod   Synced
2023-06-23T16:49:55+02:00 rbac.authorization.k8s.io ClusterRoleBinding   default  postgres-operator  Synced
2023-06-23T16:49:57+02:00          Service         default  postgres-operator  Synced Healthy
2023-06-23T16:50:00+02:00 apps            Deployment    default  postgres-operator  Synced Progressing
2023-06-23T16:50:01+02:00 acid.zalan.do    postgresql     default  acid-test-cluster  OutOfSync      postgresql.acid.zalan.do/acid-test-cluster con
figured. Warning: resource postgresqls/acid-test-cluster is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on
resources created declaratively by either create --save-config or apply. The missing annotation will be patched automatically.
2023-06-23T16:50:01+02:00          ConfigMap        default  postgres-operator  Synced      configmap/postgres-operator configured. Warnin
g: resource configmaps/postgres-operator is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources creat
ed declaratively by either create --save-config or apply. The missing annotation will be patched automatically.
2023-06-23T16:50:01+02:00 rbac.authorization.k8s.io ClusterRole   default  postgres-pod   Running Synced      clusterrole.rbac.authorization.k8s.io/postgres
-pod reconciled. reconciliation required update. clusterrole.rbac.authorization.k8s.io/postgres-pod configured. Warning: resource clusterroles/postgres-pod is missing the kubectl.kubernetes.
io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created declaratively by either create --save-config or apply. The missing ann
otation will be patched automatically.
2023-06-23T16:50:01+02:00 rbac.authorization.k8s.io ClusterRole   default  postgres-operator  Running Synced      clusterrole.rbac.authorization.k8s.io/postgres
-operator reconciled. reconciliation required update. clusterrole.rbac.authorization.k8s.io/postgres-operator configured. Warning: resource clusterroles/postgres-operator is missing the kube
ctl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created declaratively by either create --save-config or apply.
The missing annotation will be patched automatically.
2023-06-23T16:50:01+02:00 rbac.authorization.k8s.io ClusterRoleBinding   default  postgres-operator  Running Synced      clusterrolebinding.rbac.authorization.k8s.io/p
ostgres-operator reconciled. reconciliation required update. clusterrolebinding.rbac.authorization.k8s.io/postgres-operator configured. Warning: resource clusterrolebindings/postgres-operato
r is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created declaratively by either create --sa
ve-config or apply. The missing annotation will be patched automatically.
2023-06-23T16:50:01+02:00          ServiceAccount   default  postgres-operator  Synced      serviceaccount/postgres-operator configured. W
arning: resource serviceaccounts/postgres-operator is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resou
rces created declaratively by either create --save-config or apply. The missing annotation will be patched automatically.
2023-06-23T16:50:01+02:00          Service         default  postgres-operator  Synced Healthy      service/postgres-operator configured. Warning:
resource services/postgres-operator is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created d
eclaratively by either create --save-config or apply. The missing annotation will be patched automatically.
2023-06-23T16:50:01+02:00 apps            Deployment    default  postgres-operator  Synced Progressing      deployment.apps/postgres-operator configured.
Warning: resource deployments/postgres-operator is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resource
s created declaratively by either create --save-config or apply. The missing annotation will be patched automatically.
2023-06-23T16:50:02+02:00          Deployment     default  postgres-operator  Synced Healthy      deployment.apps/postgres-operator configured. Warning: resource deploym
ents/postgres-operator is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created declaratively b
y either create --save-config or apply. The missing annotation will be patched automatically.
2023-06-23T16:50:02+02:00 acid.zalan.do    postgresql     default  acid-test-cluster  Synced      postgresql.acid.zalan.do/acid-test-cluster configured. Warning: resour
ce postgresqls/acid-test-cluster is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created declar
atively by either create --save-config or apply. The missing annotation will be patched automatically.
```

ArgoCD - Lancement de la synchronisation

```
Name: argocd/postgres-operator
Project: default
Server: https://kubernetes.default.svc
Namespace: default
URL: https://10.109.47.144/applications/postgres-operator
Repo: https://github.com/simonnebaz/postgres-operator.git
Target: poc-argocd
Path: manifests
SyncWindow: Sync Allowed
Sync Policy: <none>
Sync Status: Synced to poc-argocd (98ad8c7)
Health Status: Healthy

Operation: Sync
Sync Revision: 98ad8c7aed92c6430bc36ab23228a2d57c2715d
Phase: Succeeded
Start: 2023-06-23 16:49:21 +0200 CEST
Finished: 2023-06-23 16:50:00 +0200 CEST
Duration: 39s
Message: successfully synced (all tasks run)

GROUP          KIND      NAMESPACE NAME STATUS HEALTH HOOK MESSAGE
ServiceAccount  default  postgres-operator Synced   serviceaccount/postgres-operator configured. Warning: resource serviceaccounts/postgres-o
operator is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created declaratively by either creat
e --save-config or apply. The missing annotation will be patched automatically.
ConfigMap        default  postgres-operator Synced   configmap/postgres-operator configured. Warning: resource configmaps/postgres-operator is
missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created declaratively by either --save-c
onfig or apply. The missing annotation will be patched automatically.
rbac.authorization.k8s.io ClusterRole    default  postgres-pod   Running Synced   clusterrole.rbac.authorization.k8s.io/postgres-pod reconciled. reconciliation required up
date. clusterrole.rbac.authorization.k8s.io/postgres-pod configured. Warning: resource clusterroles/postgres-pod is missing the kubectl.kubernetes.io/last-applied-configuration annotation wh
ich is required by apply. apply should only be used on resources created declaratively by either create --save-config or apply. The missing annotation will be patched automatically.
rbac.authorization.k8s.io ClusterRole    default  postgres-operator Running Synced   clusterrole.rbac.authorization.k8s.io/postgres-operator reconciled. reconciliation requir
ed update. clusterrole.rbac.authorization.k8s.io/postgres-operator configured. Warning: resource clusterroles/postgres-operator is missing the kubectl.kubernetes.io/last-applied-configuration
annotation which is required by apply. apply should only be used on resources created declaratively by either --save-config or apply. The missing annotation will be patched automatically.
rbac.authorization.k8s.io ClusterRoleBinding default  postgres-operator Running Synced   clusterrolebinding.rbac.authorization.k8s.io/postgres-operator reconciled. reconciliation
required update. clusterrolebinding.rbac.authorization.k8s.io/postgres-operator configured. Warning: resource clusterrolebindings/postgres-operator is missing the kubectl.kubernetes.io/last
-applied-configuration annotation which is required by apply. apply should only be used on resources created declaratively by either create --save-config or apply. The missing annotation
will be patched automatically.
Service         default  postgres-operator Synced   Healthy   service/postgres-operator configured. Warning: resource services/postgres-operator is mis
sing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created declaratively by either create --save-confi
g or apply. The missing annotation will be patched automatically.
apps            Deployment default  postgres-operator Synced   Healthy   deployment.apps/postgres-operator configured. Warning: resource deployments/postgres-oper
ator is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created declaratively by either create -
--save-config or apply. The missing annotation will be patched automatically.
acid.zalan.do   postgres     default  acid-test-cluster Synced   postgresql.acid.zalan.do/acid-test-cluster configured. Warning: resource postgresql/acid
-test-cluster is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by apply. apply should only be used on resources created declaratively by either
create --save-config or apply. The missing annotation will be patched automatically.
rbac.authorization.k8s.io ClusterRole    default  postgres-pod   Synced
rbac.authorization.k8s.io ClusterRole    default  postgres-operator Synced
rbac.authorization.k8s.io ClusterRoleBinding default  postgres-operator Synced
```

Installation d'argo-rollouts

```
$ kubectl create namespace argo-rollouts
namespace/argo-rollouts created
$ kubectl apply -n argo-rollouts -f https://github.com/argoproj/argo-rollouts/releases/latest/download/install.yaml
customresourcedefinition.apiextensions.k8s.io/analysisruns.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/analysistemplates.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/clusteranalystemplates.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/experiments.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/rollouts.argoproj.io created
serviceaccount/argo-rollouts created
clusterrole.rbac.authorization.k8s.io/argo-rollouts created
clusterrole.rbac.authorization.k8s.io/argo-rollouts-aggregate-to-admin created
clusterrole.rbac.authorization.k8s.io/argo-rollouts-aggregate-to-edit created
clusterrole.rbac.authorization.k8s.io/argo-rollouts-aggregate-to-view created
clusterrolebinding.rbac.authorization.k8s.io/argo-rollouts created
configmap/argo-rollouts-config created
secret/argo-rollouts-notification-secret created
service/argo-rollouts-metrics created
deployment.apps/argo-rollouts created
```

Utilisation d'Argo Rollouts

- Dans le cadre du POC, il n'a pas été nécessaire d'utiliser Argo Rollouts.
- Argo CD a été suffisant pour "pusher" la nouvelle image Docker **spilo**

Réplicaset

L'opérateur PostgreSQL met en place des réplicats set :

```
$ kubectl get rs
NAME                DESIRED  CURRENT  READY   AGE
postgres-operator-fcbd7cc96  1        1        1      101d
postgres-operator-ui-5579cc7779  1        1        1      101d
```

Détail des réplicaset

La commande suivante permet d'obtenir plus d'informations sur le réplicaset **postgres-operator** :

```
$ kubectl describe rs/postgres-operator-fcbd7cc96
Name:          postgres-operator-fcbd7cc96
Namespace:     default
Selector:      name=postgres-operator,pod-template-hash=fcbd7cc96
Labels:        name=postgres-operator
               pod-template-hash=fcbd7cc96
Annotations:   deployment.kubernetes.io/desired-replicas: 1
               deployment.kubernetes.io/max-replicas: 1
               deployment.kubernetes.io/revision: 1
Controlled By: Deployment/postgres-operator
Replicas:      1 current / 1 desired
Pods Status:   1 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:      name=postgres-operator
               pod-template-hash=fcbd7cc96
  Service Account: postgres-operator
  Containers:
    postgres-operator:
      Image:      registry.opensource.zalan.do/acid/postgres-operator:v1.9.0
      Port:       <none>
      Host Port: <none>
      Limits:
        cpu:      500m
        memory:  500Mi
      Requests:
        cpu:      100m
        memory:  250Mi
  Environment:
    CONFIG_MAP_NAME:  postgres-operator
  Mounts:
  Volumes:
  Events:      <none>
```

StatefulSets

- La définition officielle des StatefulSets est disponible à l'URL suivante⁷.
- Les Pods portants l'image Spilo sont déployés en se basant sur des statefulsets :

```
$ kubectl get statefulsets  
NAME           READY   AGE  
acid-test-cluster  2/2    88d
```

7.

<https://kubernetes.io/docs/concepts/workloads/controllers/statefulset/>

Détails des statefulSets

```
# kubectl describe statefulset.acid-test-cluster
Name:           acid-test-cluster
Namespace:      default
CreationTimestamp: Wed, 29 Mar 2023 16:51:13 +0200
Labels:         application=spilo
                cluster-name=acid-test-cluster
                team=acid
Annotations:    
Replicas:      2 desired | 2 total
Pod难得:       2 Running / 0 Waiting / 0 Failed
Pod Template:
  Labels:  application=spilo
          cluster-name=acid-test-cluster
          team=acid
  Service Account: postgres-pod
  Init Containers:
    data:
      Image:        postgres
      Port:        <none>
      Host Port:   <none>
      Container ID: <none>
      Environment: <none>
      Mounts:      <none>
      Limits:      <none>
      Requests:    <none>
      CPU:         100m
      Memory:     100Mi
  Requests:
    CPU:        100
    Memory:    100Mi
  Environment:
    SCOPES:      acid-test-cluster
    PGPORT:     /home/postgres/pgdata/postgres
    PGD_IP:      (viinstatistics.podIP)
    PGD_NAMESPACE: (viinstatistics.namespace)
    PGD_SERVICE:  postgres
    KUBERNETES_SCOPE_LABEL: cluster-name
    KUBERNETES_ROLE_LABEL: spilo-role
    PGADMIN_SUPERUSER: <use the key 'password' in secret 'postgres.acid-test-cluster.credentials.postgresql.acid.ralan.eu'> optional: false
    PGADMIN_STANDBY: <use the key 'password' in secret 'standby.acid-test-cluster.credentials.postgresql.acid.ralan.eu'> optional: false
    PGADMIN_STANDBY2: <use the key 'password' in secret 'standby2.acid-test-cluster.credentials.postgresql.acid.ralan.eu'> optional: false
    PAM_SPAUTH:   https://infra.example.com/auth/tokensinfo/token_realm/employees
    PAM_SPAUTH2:  https://infra.example.com/auth/tokensinfo/token_realm/employees
    EMAIL_MAL_PATH_COMPAT: true
    EMAIL_MAL_PATH_COMPAT2: true
    PGVERSION:   "15"
    SPILLOUTPUTS:  ["application","spilo"]
    SPILLOUTPUTS2:  ["postgresql"]
    SPILLOUTPUTS3:  ["postgresql","parameters":{"shared_buffers":"128M"}]
    Volumes:
      pgdata:          persistentVolumeClaim_acid_ralan_eu: true
      mount0:          persistentVolumeClaim_acid_ralan_eu: true
      /dev/zero:        emptyDir (a temporary directory that shares a pod's lifetime)
      /home/postgres/pgdata: emptyDir (rw)
      /home/postgres/pgdata: emptyDir (rw)
      /opt/tmpdir:      emptyDir (rw)
  Volumes:
    pgdata:
      Type:      EmptyDir (a temporary directory that shares a pod's lifetime)
      Medium:    Memory
      SizeLimit: <none>
      options:  <none>
    mount0:
      Type:      EmptyDir (a temporary directory that shares a pod's lifetime)
      Medium:    Memory
      SizeLimit: <none>
      options:  <none>
  Volume Claims:
    Name:          pgdata
    StorageClass: standard
    Labels:        application=spilo
                  cluster-name=acid-test-cluster
                  team=acid
    Annotations:  
    Capacity:     5Gi
    Storage Mode: [ReadWriteOnce]
    Events:       <none>
```

Détails des statefulSets

- Comme indiqué dans les détails du statefulset, la politique de mise à jour est **OnDelete**.
- D'après la documentation officielle, cela signifie que la mise à jour des pods passe par une suppression des pods.

Traitement des statefulsets par Argo

- A priori, le traitement des statefulsets peut avoir des limitations⁸
- Cette limitation a été corrigée

8.

<https://argo-cd.readthedocs.io/en/stable/faq/#why-is-my-application-stuck-in-progressing-state>

Mise à jour des images Spilo - Upgrade mineur

- Le script de création du cluster mis à disposition par le dépôt de l'opérateur PostgreSQL est : **manifests/complete-postgres-manifest.yaml**
- Il va être copié dans un répertoire dédié (**deployment**) pour gérer sa synchronisation avec ArgoCD

```
argocd app create postgres-deployment \
--repo https://github.com/simonelbaz/postgres-operator.git \
--path deployment \
--dest-server https://kubernetes.default.svc \
--revision poc-argocd \
--dest-namespace default
```

- Le script est modifié avec la nouvelle version de l'image spilo :

```
$ cat deployment/complete-postgres-manifest.yaml | grep dockerImage
dockerImage: ghcr.io/zalando/spilo-15:3.0-p1
```

Mise à jour des images Spilo - Upgrade mineur

```
$ argocd app sync postgres-deployment
TIMESTAMP          GROUP           KIND    NAMESPACE      NAME    STATUS   HOOK   MESSAGE
2023-06-26T18:44:36+02:00  acid.zalan.do  postgresql  default  acid-test-cluster  OutOfSync
2023-06-26T18:45:07+02:00  acid.zalan.do  postgresql  default  acid-test-cluster  OutOfSync
postgresql.acid.zalan.do/acid-test-cluster configured
2023-06-26T18:45:09+02:00  acid.zalan.do  postgresql  default  acid-test-cluster  Synced
postgresql.acid.zalan.do/acid-test-cluster configured

Name:          argocd/postgres-deployment
Project:       default
Server:        https://kubernetes.default.svc
Namespace:     default
URL:          https://10.109.47.144/applications/postgres-deployment
Repo:          https://github.com/simonelbaz/postgres-operator.git
Target:        poc-argocd
Path:          deployment
SyncWindow:    Sync Allowed
Sync Policy:   <none>
Sync Status:   Synced to poc-argocd (cfb2094)
Health Status: Healthy

Operation:     Sync
Sync Revision: cfb20940760cf96ced4a5872f161c280457c9784
Phase:         Succeeded
Start:         2023-06-26 18:44:34 +0200 CEST
Finished:      2023-06-26 18:45:06 +0200 CEST
Duration:      32s
Message:       successfully synced (all tasks run)

GROUP          KIND    NAMESPACE  NAME      STATUS  HEALTH  HOOK   MESSAGE
acid.zalan.do  postgresql  default  acid-test-cluster  Synced
                                                               postgresql.acid.zalan.do/acid-test-cluster configured
```

Accès client à la base de données

- Le lien suivant décrit comment se logger à la base de données PostgreSQL⁹
- Récupération du pod master :

```
# get name of master pod of acid-minimal-cluster
export PGMMASTER=$(kubectl get pods \
-o jsonpath=.items..metadata.name \
-l application=spilo,cluster-name=acid-minimal-cluster,spilo-role=master \
-n default)

# set up port forward
kubectl port-forward $PGMASTER 6432:5432 -n default
```

9.

<https://github.com/zalando/postgres-operator/blob/master/docs/user.md#connect-to-postgresql>

Accès client à la base de données

- Connexion depuis un tunnel SSH :

```
export PGPASSWORD=$(kubectl get secret \
postgres.acid-minimal-cluster.credentials.postgresql.acid.zalan.do \
-o 'jsonpath=.data.password' | base64 -d)
export PGSSLMODE=require
psql -U postgres -h localhost -p 6432
```

Point-in-time recovery PITR

Vidéo partagée 7-Postgres_operator_pitr.mkv

Failover - Changement du nombre de pods

Vidéo partagée *6-Postgres_operator_failover_podnumberchange.mkv*

Mise à jour mineure de PostgreSQL

Vidéo partagée *6-Postgres_operator_minorupgrade.mkv*

Mise à jour majeure de PostgreSQL

Vidéo partagée *6-Postgres_operator_major_upgrade.mkv*

Supervision

Supervision du cluster k8s - Prometheus

- Le cluster k8s nécessite une supervision.
- Prometheus est une solution de supervision open source répandue
- Elle possède un opérateur **kube-prometheus** dont le déploiement est décrit depuis le site Github

Déploiement de kube-prometheus

- Le lien suivant décrit les commandes de déploiement¹⁰

10.

<https://github.com/prometheus-operator/kube-prometheus#quickstart>

Déploiement de kube-prometheus

- Les commandes suivantes décrivent le déploiement de l'opérateur **kube-prometheus**

```
$ git clone https://github.com/prometheus-operator/kube-prometheus.git
$ kubectl apply --server-side -f manifests/setup
customresourcedefinition.apiextensions.k8s.io/alertmanagerconfigs.monitoring.coreos.com serverside-applied
...
$ kubectl wait \
  --for condition=Established \
  --all CustomResourceDefinition \
  --namespace=monitoring
customresourcedefinition.apiextensions.k8s.io/alertmanagerconfigs.monitoring.coreos.com condition met
customresourcedefinition.apiextensions.k8s.io/alertmanagers.monitoring.coreos.com condition met
customresourcedefinition.apiextensions.k8s.io/analysisruns.argoproj.io condition met
...
$ kubectl apply -f manifests/
...
```

- la dernière commande s'est terminée en erreur et a nécessité un double lancement

DaemonSets - Prometheus

- Le déploiement des nodes exporter de Prometheus est instancié sous forme de DaemonSets
- Liste des daemonsets

```
$ kubectl get daemonsets -n monitoring
NAME        DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
node-exporter   10        10        10      10           10          kubernetes.io/os=linux   110m
```

Accès à l'UI de Prometheus

- le lien suivant décrit le mode opératoire pour accéder à l'UI de Prometheus¹¹
- Depuis le nœud control plane, ouvrir un port forward :

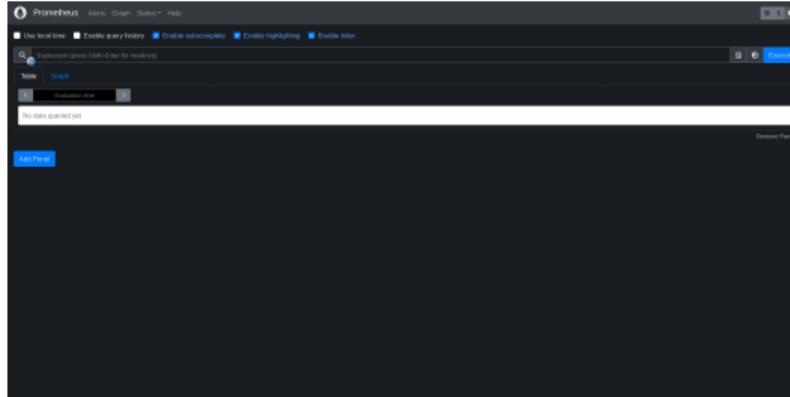
```
$ kubectl --namespace monitoring port-forward svc/prometheus-k8s 9090
Forwarding from 127.0.0.1:9090 -> 9090
Forwarding from [::1]:9090 -> 9090
```
- Pour accéder à l'interface web de l'opérateur PostgreSQL depuis le PC de l'utilisateur, il est possible de passer par une redirection SSH :

```
ssh -L 9095:10.106.57.137:9090 dgfip-k8s
```
- Lancer le navigateur pour accéder à l'URL `http://localhost:9095/`

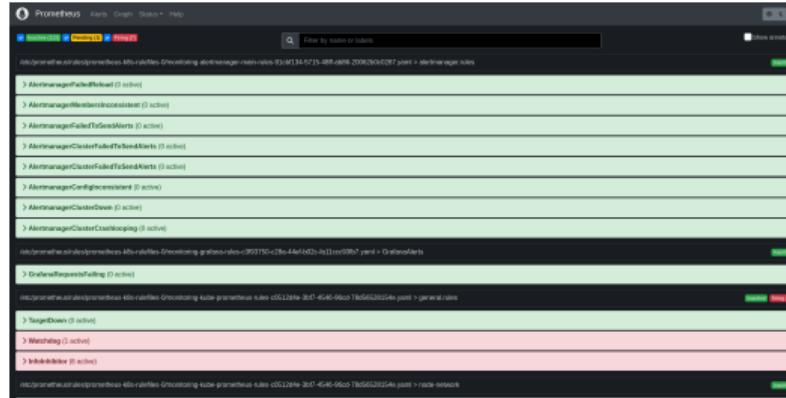
11.

<https://github.com/prometheus-operator/kube-prometheus/blob/main/docs/access-ui.md>

Ecran d'accueil de l'UI de Prometheus



Ecran des alertes de l'UI de Prometheus



Flags de la ligne de commandes - Prometheus

Flag	Value	Description
--alertmanager-connection-capacity	10000	
--alertmanager-dlstats		Whether to use the alert manager's dlstats endpoint.
--config_file		
--enable-feature		
--log-format	logfmt	
--log-level	info	
--query-blocksize-data	5m	
--query-max-concurrent	20	
--query-max-samples	5000000	
--query-timestamp	2m	
--rules-alert-for-grace-period	10m	
--rules-alert-for-warning-tolerance	1h	
--rules-alert-reload-delay	1h	
--storage-cdp-allow-timestamp	true	
--storage-cf-expiry-reload-interval	5s	
--storage-disk-dump-tolerance	2ms	
--storage-disk-err-tolerance	false	
--storage-agent-cash	data-agent	
--storage-agent-reload-interval-new-time	0s	
--storage-agent-reload-interval-new-time-base	0s	
--storage-agent-wal-compression	no	

Rules - Prometheus

The screenshot shows the Prometheus Rules configuration page. There are three entries listed under the heading "alertmanager.rules". Each entry includes a detailed description of the alert rule, its state (green), and its last evaluation time.

Rule	State	Error	Last Evaluation	Evaluation Time
<pre>alert AlertmanagerRuleFailed expr: sum_over_time(alertmanager_cluster_size[1m]) <= total_size * 'warning' * 0.95 for: 1m labels: severity: critical annotations: description: Configuration has failed to load for [[Labels]] namespace [[0]]. Below is the configuration that failed: value: apiVersion: v1 kind: AlertmanagerConfiguration metadata: name: alertmanager spec: receiver: 'receiver' route: cluster: cluster_name: 'cluster' cluster_size: 1 selector: match_labels: receiver: 'receiver'</pre>	green		9.189s ago	0.259ms
<pre>alert AlertmanagerMemberNotInCluster expr: sum_over_time(alertmanager_member_in_cluster[1m]) <= total_size * 'warning' * 0.95 for: 1m labels: severity: critical annotations: description: Alertmanager [[Labels]] namespace [[0]] has only found [[1]] members of the [[Labels]] pod(s) cluster. value: apiVersion: v1 kind: AlertmanagerConfiguration metadata: name: alertmanager spec: receiver: 'receiver' route: cluster: cluster_name: 'cluster' selector: match_labels: receiver: 'receiver'</pre>	green		9.189s ago	0.190ms
<pre>alert AlertmanagerFederatedServices expr: sum_over_time(alertmanager_federated_services[1m]) <= total_size * 'warning' * 0.95 for: 1m labels: severity: warning</pre>	green		9.189s ago	0.529ms

Runtime - Prometheus

The screenshot shows the Prometheus 'Runtime Information' page. It includes sections for 'Runtime Information', 'Build Information', and 'Alertmanagers'.

Runtime Information

Key	Value
Start time	Tue, 27 Jun 2023 20:38:15 GMT
Working directory	/var/lib/prometheus
Configuration reload	Successful
Last successful configuration reload	2023-06-27T20:30:27Z
HTTP port range	0
Graceful shutdown timeout	30s
GOMAXPROCS	2
GOMAXBITS	64
GOGC	
GODEBUG	
Storage retention	1d

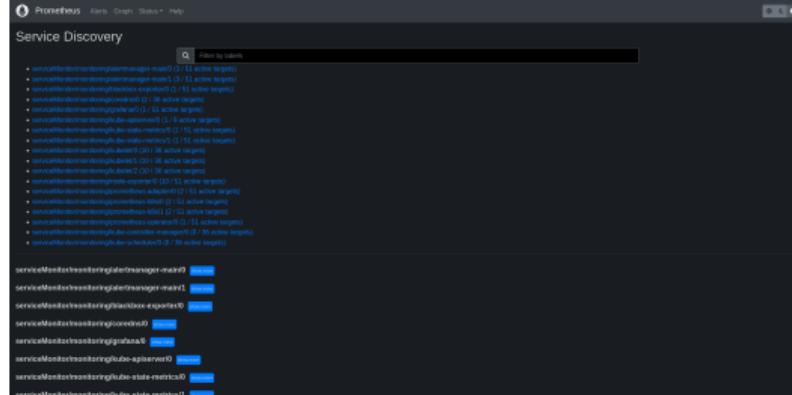
Build Information

Key	Value
Version	2.45.0
Revision	6a777c7e096944a9f309b95a187e7ff73271448
Branch	HEAD
BuildUser	rancher@90033964657
BuildDate	2023-06-27 16:09:49
GoVersion	go1.20.5

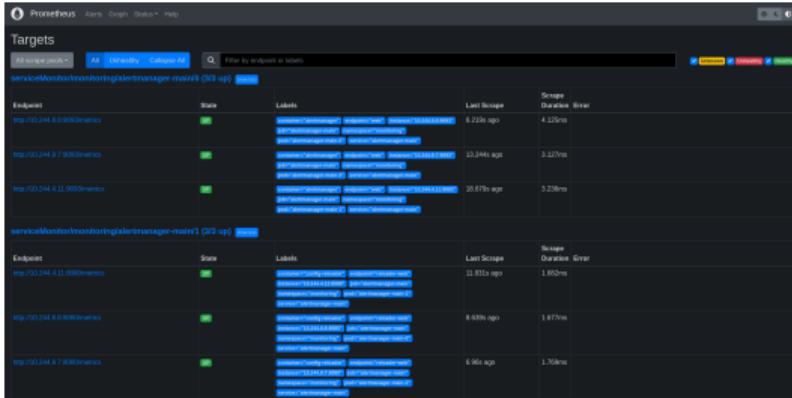
Alertmanagers

Endpoint	URL
localhost	http://127.0.0.1:13900/api/v1/alerts
localhost	http://127.0.0.1:13900/api/v1/alerts

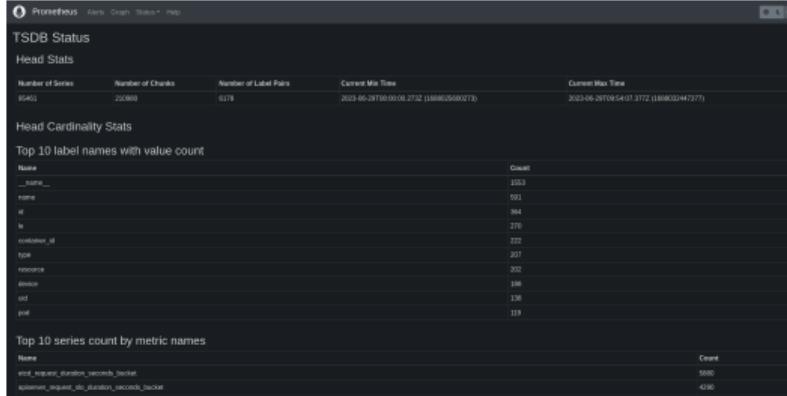
Service Discovery - Prometheus



Targets - Prometheus



TSDB Status - Prometheus



Déploiement du node exporter de PostgreSQL

- Le précédent déploiement permet de moniterer le cluster k8s
- Il est maintenant temps de moniterer la base de données PostgreSQL
- La communauté Prometheus met à disposition l'exporter PostgreSQL à l'URL suivante :
https://github.com/prometheus-community/postgres_exporter
- Il est possible de déployer l'exporter postgres sur un pod de type **sidecar**
- Par simplicité du POC, l'image Docker mise à disposition par la communauté a été employée

Déploiement du node exporter de PostgreSQL

- Les commandes de déploiement de l'exporter PostgreSQL sont décrites dans l'URL suivante : https://github.com/prometheus-community/postgres_exporter#quick-start

Lancement de l'exporter PostgreSQL

Dans un 1^{er} temps, on met en place le classique port forward depuis le control plane :

```
$ export PGMASTER=$(kubectl get pods \
-o jsonpath=.items..metadata.name \
-l application=spilo,cluster-name=acid-test-cluster,spilo-role=master \
-n default)
$ kubectl port-forward $PGMASTER 6432:5432 -n default
Forwarding from 127.0.0.1:6432 -> 5432
Forwarding from [::1]:6432 -> 5432
```

Dans un 2^e temps, le Docker est lancé depuis la VM dédiée :

```
sudo docker run --net=host -p 9187:9187 \
-e DATA_SOURCE_NAME="postgresql://postgres:*****@localhost:6432/postgres?sslmode=require" \
-v /home/linagora/postgres_exporter.yml:/home/linagora/postgres_exporter.yml \
quay.io/prometheuscommunity/postgres-exporter \
--config.file="/home/linagora/postgres_exporter.yml" \
--log.level="debug"
```

Accès à l'UI de l'exporter Postgres

- Ouvrir un tunnel SSH vers le Docker exporter Postgres :
`ssh -L 9187:localhost:9187 dgfip-prometheus`
- Lancer le navigateur pour accéder à l'URL `http://localhost:9187/`

UI de l'exporter Postgres



Extrait de métriques remontées par l'exporter PostgreSQL

```
# HELP pg_statio_user_tables_heap_blocks_read Number of disk blocks read from this table
# TYPE pg_statio_user_tables_heap_blocks_read counter
pg_statio_user_tables_heap_blocks_readdatname="postgres",relname="job",schemaname="cron" 0
pg_statio_user_tables_heap_blocks_readdatname="postgres",relname="job_run_details",schemaname="cron" 0
pg_statio_user_tables_heap_blocks_readdatname="postgres",relname="postgres_log",schemaname="public" 0
# HELP pg_statio_user_tables_idx_blocks_hit Number of buffer hits in all indexes on this table
# TYPE pg_statio_user_tables_idx_blocks_hit counter
pg_statio_user_tables_idx_blocks_hittestname="postgres",relname="job",schemaname="cron" 0
pg_statio_user_tables_idx_blocks_hittestname="postgres",relname="job_run_details",schemaname="cron" 0
pg_statio_user_tables_idx_blocks_hittestname="postgres",relname="postgres_log",schemaname="public" 0
# HELP pg_statio_user_tables_idx_blocks_read Number of disk blocks read from all indexes on this table
# TYPE pg_statio_user_tables_idx_blocks_read counter
pg_statio_user_tables_idx_blocks_readdatname="postgres",relname="job",schemaname="cron" 0
pg_statio_user_tables_idx_blocks_readdatname="postgres",relname="job_run_details",schemaname="cron" 0
pg_statio_user_tables_idx_blocks_readdatname="postgres",relname="postgres_log",schemaname="public" 0
# HELP pg_statio_user_tables_tidx_blocks_hit Number of buffer hits in this table's TOAST table indexes (if any)
# TYPE pg_statio_user_tables_tidx_blocks_hit counter
pg_statio_user_tables_tidx_blocks_hittestname="postgres",relname="job",schemaname="cron" 0
pg_statio_user_tables_tidx_blocks_hittestname="postgres",relname="job_run_details",schemaname="cron" 0
pg_statio_user_tables_tidx_blocks_hittestname="postgres",relname="postgres_log",schemaname="public" 0
# HELP pg_statio_user_tables_tidx_blocks_read Number of disk blocks read from this table's TOAST table indexes (if any)
# TYPE pg_statio_user_tables_tidx_blocks_read counter
pg_statio_user_tables_tidx_blocks_readdatname="postgres",relname="job",schemaname="cron" 0
pg_statio_user_tables_tidx_blocks_readdatname="postgres",relname="job_run_details",schemaname="cron" 0
pg_statio_user_tables_tidx_blocks_readdatname="postgres",relname="postgres_log",schemaname="public" 0
```

Sommaire

Sommaire

Architecture

Installation

Stockage

Déploiement PostgreSQL

Stockage S3 - Minio

Sécurité - Support d'OpenShift

Déploiement continu



Fin du document

